

An Observational Benchmark for Layer 2 Finality

What has been built, why each piece exists, how the parts fit together, and what every file in the repository is for.

PROJECT	50081 — Enhancing Scalability and Interoperability in Decentralized Layer 2 Blockchain Networks
TRACK	Scalability — benchmarking L2 approaches on throughput, latency and cost-efficiency
PROGRAMME	Mitacs Globalink 2026
SUPERVISOR	Prof. Sara Rouhani, TCDT Lab
INTERN	Ujala Kiran
APPROACH	Observational — measure public networks, host nothing
STATUS	24 of 38 tasks complete. Phases A–D done; minimum viable deliverable reached
DATE	10 August 2026

WHAT THIS DOCUMENT IS

The implementation handbook says *what* to build, task by task. This document records *what was actually built*, what was learned in the process, and why the code looks the way it does. Where a measurement contradicted an expectation, that is written down here rather than quietly fixed — several of those contradictions are the most useful findings so far.

1 · What the project measures

The framework submits a transaction to a Layer 2 network operated by somebody else, then determines three different moments at which that transaction could reasonably be called *final* — and what it cost to reach each one.

LEVEL	FINAL WHEN	WHERE THE TIMESTAMP COMES FROM
t₁ Full trust	The L2 sequencer accepts it	The L2's own RPC — a transaction receipt appears
t₂ Partial trust	Its batch is posted to Ethereum	An Ethereum L1 transaction sent by the rollup operator
t₃ Trustless	Its batch is proven, or its challenge window closes	A second Ethereum L1 transaction, later

Two of the three timestamps come from Ethereum, not from the rollup. That single observation is why the project needs no hosted infrastructure: the evidence for partial-trust and trustless finality is public on the L1, sent by the rollup operator, and readable by anyone. It is also what makes every result independently checkable — each row of output carries the Ethereum transaction hashes a reviewer can open for themselves.

Why three levels rather than one

Published L2 benchmarks report "latency" as a single number, which almost always means t_1 — the sequencer said yes. That number is excellent for every rollup and says nothing about the security assumption behind it. Separating the three levels is what allows the comparison this project exists to make: a ZK rollup and an optimistic rollup look similar at full trust and behave very differently once you stop trusting the sequencer.

2 · Where the project stands

PHASE	PURPOSE	TASKS	STATUS
A	Unblock	A1–A5	Complete
B	Make it real	B1–B5	Complete
C	The contribution	C1–C6	Complete
D	Make it defensible	D1–D5	Complete — minimum viable deliverable
E	The comparison	E1–E3	Not started — OP Sepolia funded and ready
F	Data collection	F1–F4	Not started
G	Analysis	G1–G5	Not started
H	Write and deliver	H1–H5	Not started

Results obtained so far

zkSync Era Sepolia, 109 transactions, all successful, all settled to Ethereum Sepolia in batch 21623.

LEVEL	MEDIAN	P95	MIN	MAX
Full trust (t_1)	18.04 s	19.19 s	0.60 s	20.16 s
Partial trust (t_2)	25.47 m	28.30 m	19.48 m	30.42 m
Trustless (t_3)	36.67 m	39.50 m	30.68 m	41.62 m

Broken down by stage, the medians are: 0.30 minutes for the sequencer to accept, **25.17 minutes waiting for the batch to be posted**, and 11.20 minutes from commit to proof. The dominant cost of trustless finality on this rollup is not proving — it is waiting for a batch to seal.

Cost per transaction: **\$0.00285965** at an assumed ETH price of \$3,800, derived from the actual L1 receipts (commit 0.000495 ETH, proof 0.000164 ETH) divided by the 875 transactions sharing the batch.


THE CAVEAT THAT GOVERNS EVERY NUMBER ABOVE

109 transactions, but **one batch**. One commit transaction, one proof transaction, submitted inside a single eleven-minute window. The tell is that the t_2 and t_3 spreads are identical — 10.95 minutes each — because t_2 and t_3 are the same two constants for all 109 rows and the entire spread comes from our own submission times differing.

Settlement itself therefore has **$n = 1$** . These medians demonstrate that the pipeline works end to end; they are not yet a characterisation of the network. Task F2 — five repetitions, at least an hour apart — is what turns them into one.

3 · How the system flows

The single most important structural decision is that submission and resolution are **two separate programs**. Trustless finality can take an hour. A program that waits synchronously ties up the machine, dies when the laptop sleeps, and loses everything it had collected.



Why this shape

- **t₀ is captured immediately before the broadcast**, after signing and gas estimation. Anything expensive between the two would be silently attributed to the network.
- **The file is the state.** The resolve pass recomputes what is outstanding from disk on every pass and holds nothing between passes, so it can be interrupted at any point and lose nothing.
- **JSON Lines, not a JSON array.** An interrupted append leaves a readable file; a half-written array does not.
- **Only absolute timestamps are stored.** Latencies are computed at export time. An arithmetic mistake becomes a re-export rather than a re-run of the experiment.

4 · Repository layout

```
bench/
  core/          submit, resolve, metrics — network-independent
  adapters/      one file per rollup — the only per-network code
  configs/       one YAML per concern; a run is described entirely by these
  abis/          committed contract ABIs
  contracts/     the workload token and its vendored dependencies
  results/       gitignored JSONL output
  export/        gitignored CSV deliverables
  docs/          this handbook and its build script
```

The separation that matters is `core/` against `adapters/`. Everything in `core/` works in terms of a `Settlement` record and never asks which rollup it is talking to. Adding a second architecture means adding one file to `adapters/` and changing nothing else — which is what makes the ZK-versus-optimistic comparison achievable rather than aspirational.

5 · Every file, and what it is for

Entry points — the commands you actually run

`bench/create_wallet.py` A2

Generates the single test account used across every network and writes the private key to `.env`, the address to committed config. One account everywhere keeps records simple, and addresses are compatible across all EVM chains.

```
main()
```

`bench/check_funds.py` A3 · A4

Answers "where do I actually have money" without opening a wallet, and doubles as the endpoint verification A4 requires. Reports chain ID mismatches and — importantly — **stalled endpoints**: an RPC that answers correctly while serving state frozen weeks in the past. Records balances and faucet provenance into `accounts.yaml` for the report's experimental setup.

```
probe · render · next_actions · print_faucets · record
```

`bench/send_one.py` A5

A deliberately unabridged script that sends one transaction and times it. It exists so that when the framework misbehaves later, there is a known-good reference proving that keys, funds, endpoint and chain ID all work together. `--dry-run` signs locally and needs no funds at all.

```
build · wait_for_receipt · main
```

`bench/deploy_token.py` B1

Compiles `BenchToken.sol` with a pinned solc and deploys it, recording the address into `workloads.yaml` by a targeted line edit that preserves the file's comments. `--dry-run` asks the node whether it will accept the bytecode at all, which is how the question "does this rollup need its own compiler" gets answered by measurement instead of assumption.

```
compile_token · record_address · main
```

bench/check_workloads.py B1 DONE-WHEN

Builds every workload against a live node and asks for a gas estimate without broadcasting. Estimation is a real simulation against real state, so it catches a corrupt selector, a wrong token address and an unfunded account — which is why it is the acceptance test rather than a local encoding check.

describe · check · run_network · main

bench/submit_run.py B2 · B3 · B4 · B5

The submit pass. Builds a batch of one workload, numbers it consecutively from a single fetched nonce, broadcasts in index order, polls every receipt round-robin, and appends one JSON object per transaction. Exits as soon as t_1 is known — it never waits for settlement.

default_run_id · main

bench/resolve_run.py C2-C6

The resolve pass. Reads run files, finds rows still missing t_2 or t_3 , asks the rollup where they settled, and writes the files back. Safe to run repeatedly and safe to interrupt. `--watch N` repeats every N seconds and stops itself once nothing is outstanding.

resolve_file · one_pass · main

bench/export.py D1 · D2 · D3 · D4

Produces the three deliverable files. Computes all latencies here from absolute timestamps, fetches L1 costs once per batch rather than once per transaction, and writes an assumptions block stating the ETH rate, its source, and every declared limitation. `--eth-usd` overrides the configured rate, which is D2's acceptance test.

latencies · write_raw · batch_costs · summarise_group · main

bench/check_config.py D5 DONE-WHEN

Finds configuration keys that nothing reads. A setting that is parsed and then ignored is the worst class of benchmark bug: you change it, the run completes, the number is unchanged, and nothing tells you. It found three such keys on first run.

schema_keys · main

Core — network-independent machinery

bench/core/wallet.py A2

Loads the private key from the environment and confirms it derives the address recorded in committed config. A mismatch means `.env` and the config describe different accounts, which would make every recorded result untraceable.

load_env · load_account · configured_address · account_created · check_address

bench/core/networks.py A3 · A4 · C1

Every endpoint, chain ID, explorer URL and L1 contract address, loaded from YAML and never inline in code, so a run is described entirely by committed config. `connect_verified` refuses to proceed when an endpoint reports a different chain than the config claims — an endpoint quietly pointing at the wrong chain produces plausible results describing the wrong system.

Network · load_networks · funding_priority · connect · connect_verified

bench/core/workloads.py B1

Turns a workload definition into a transaction dict. Enforces the rule that **no selector and no calldata is ever written by hand**: selectors are derived from the committed ABI and arguments go through the ABI encoder. Gas is deliberately not set here, so that an estimation failure is reported as an estimation failure rather than becoming a hardcoded limit.

Workload · load_workloads · recipient · token_address · selector · signature · build

bench/core/submit.py B2 · B3 · B4 · B5

Signing, broadcasting, nonce assignment, receipt polling and outcome classification. Captures t_0 immediately before the broadcast. Raises on submission failure so the caller can record it with `hash` left null — **a fabricated hash is the single most damaging thing that can be done to a benchmark**, because every downstream measurement then looks plausible and means nothing.

to_hex · estimate_gas · assign_nonces · verify_nonce_advance · submit · rejected_record · resolve_t1_batch · annotate_block_times · summarise · group_failures

bench/core/records.py B2 · C5 · C6 · D1

The row schema and the JSON Lines file format. Defines the outcome vocabulary — success, reverted, timeout, rejected — and normalises rows on load so a file written before a field existed still reads as a rectangular table. Writes via a temporary file and an atomic move, so an interrupted pass cannot leave a half-written run file.

Outcome · new_record · run_path · append · normalise · load · rewrite · iter_runs

bench/core/settle.py C2 · C3 · C4 · C5

Turns L1 settlement hashes into timestamps: fetch the Ethereum transaction, read its block number, fetch that block, take its timestamp. Caches per L1 transaction, because a batch holds many of our rows and re-fetching each block once per row multiplies the request count for no new information. `outstanding()` decides what work remains.

l1_timestamp · apply_settlement · outstanding

bench/core/costs.py D2 · D3

Cost and data-availability size from L1 receipts — `gasUsed × effectiveGasPrice`, never a constant and never an estimate. Distinguishes blob-carrying transactions (type 3, priced at 131,072 bytes per blob) from calldata postings (16 gas per non-zero byte). Bytes and cost travel as separate fields because the byte count is a fact about the rollup while the cost depends on our price assumptions.

L1Cost · l1_cost · pricing · to_usd

Adapters — the only per-rollup code

bench/adapters/base.py C2 · E1

Defines `Settlement`, the one shape of answer every architecture returns, so the resolve pass never branches on which rollup it is talking to. Carries `t3_kind` — observed or derived — because a ZK rollup *witnesses* a proof while an optimistic rollup *computes* a deadline, and presenting the second as the first would be the most misleading thing in the study.

`Settlement · Adapter · SettlementUnavailable · OBSERVED · DERIVED`

bench/adapters/zksync.py C2

Resolves settlement via the L2 receipt's batch number and `zks_getL1BatchDetails`, caching per batch. Also yields the batch's own transaction count, which is the denominator the cost model needs.

`ZkSyncAdapter · batch_details · settlement`

bench/adapters/__init__.py C2 · E1

Maps a network's `family` to its adapter. Raises rather than returning a do-nothing default, because a rollup with no adapter would otherwise produce rows with no t_2 and no t_3 that look exactly like rows whose settlement has not happened yet.

`for_network`

Configuration — a run is described entirely by these

bench/configs/networks.yaml

Five chains: one L1 and four L2s. Endpoints, expected chain IDs, explorers, bridge URLs, faucet tiers, funding priority, and the L1 contract addresses two of the three timestamps are read from. Endpoints can be overridden per network from the environment (`BENCH_RPC_<KEY>`) to swap a public endpoint for a paid one.

bench/configs/workloads.yaml

Two workloads — a native transfer and an ERC-20 transfer — plus the fixed recipient and the deployed token address per network. Contains no selector and no calldata: those are derived from the ABI at runtime.

bench/configs/accounts.yaml

The test account's address, when it was created, and — recorded as it happens rather than remembered in Week 7 — which faucet or bridge funded each network. This is the report's experimental-setup section in raw form.

bench/configs/pricing.yaml

The ETH/USD rate, when it was taken, and from where. Read only at analysis time, never at collection time, so a wrong rate is a one-line fix and a re-export rather than a re-run. A dollar figure without its assumptions cannot be checked by anyone.

bench/abis/erc20.json · zksync_hyperchain.json

Committed rather than fetched at runtime, so results do not stop being reproducible because a third-party API is down or has changed its terms. The zkSync ABI carries the `BlockCommit` and `BlocksVerification` events the fallback route needs, and the batch counters mainnet observation will use.

bench/contracts/BenchToken.sol

The plainest possible OpenZeppelin ERC-20: no hooks, no pausing, no permit. The study measures what it costs a rollup to settle an ordinary token transfer, so anything unusual in this contract would appear in the gas and DA figures as though it were a property of the rollup. OpenZeppelin sources are vendored at a pinned 5.0.2 alongside it.

6 · What was learned by measuring

Each of these contradicted a reasonable expectation. They are recorded because they are the findings most likely to matter to somebody repeating this work.

The per-transaction settlement endpoint lags the per-batch one

The obvious route to settlement is `zks_getTransactionDetails`, which returns commit, prove and execute hashes for an L2 transaction. For batch 21623 it returned `ethProveTxHash: null` while `zks_getL1BatchDetails` for that same batch already carried a proof hash timestamped seven minutes after the commit. The transaction view catches up eventually, but any single pass over it **systematically undercounts trustless finality**, and nothing about the null value signals that it is stale rather than genuinely pending. Settlement is now read per batch.

The batch-posting transaction's recipient is not the rollup's contract

The address a commit transaction is sent to (`0xedf4B1C6...`) is the ValidatorTimelock — batches are submitted *through* it, and every state getter on it reverts. The contract holding batch state is a different address (`0x9a6de0f6...`), obtained from the chain itself via `zks_getMainContract`. A fallback route that queried logs on the timelock would find nothing and look like a broken query rather than a wrong address.

zkSync reverts on transfers to low addresses

The conventional burn address `0x...dEaD` is the integer 57,005, which falls inside the range zkSync reserves for system contracts. Transfers to it revert with a bare `execution reverted` and no reason. It estimates fine on Sepolia and fails only on the rollup — precisely the kind of thing that surfaces as an inexplicable Phase B failure. The recipient is now a deterministically derived address outside that range.

Gas estimates are not cost figures

On zkSync, the token deployment estimated 5,386,132 gas and used 590,512 — a ninefold overshoot. A transfer estimated 224,386 and used 79,259. On OP Sepolia the same deployment estimated 555,635 and used 550,102, which is accurate. Costs must come from receipts; an estimate-derived cost on zkSync would not be wrong by a little.

Batch size drives full-trust latency

Median t_1 was 0.61 s for a single transaction, 1.85 s for five, and 18.11 s for fifty. The first reading was that this must be an artefact of the polling interval — halving the poll cap moved the figure from 18.63 s to 17.51 s, which disproved it. It is the sequencer working through a queue of consecutive nonces from one account, and it is the effect the batch-size sweep exists to characterise.

Per-transaction data-availability size is not measurable here

A blob costs 131,072 bytes whether the rollup fills it or not, and batch 21623 carried 875 transactions of which 766 were not ours. So per-transaction DA bytes is a batch average identical across workloads, and the acceptance test for that task — DA bytes differing between a native and an ERC-20 batch — is **not reachable on this rollup at this scale**. The column is named `da_bytes_per_tx_batch_avg`, flagged `da_workload_sensitive=False`, and the limitation is stated in the export rather than hidden by a plausible-looking number. The workload difference is real and appears in L2 gas, where zkSync prices pubdata: 82,995 for a native transfer against 121,916 for an ERC-20 transfer.

Full-trust latency depends on what you mean by it

Measured t_1 exceeds the L2 block's own timestamp by a median of 17.95 s for batched submissions. The sequencer timestamps the block near submission but does not expose the receipt for some seconds afterwards. So t_1 as measured means "when the sequencer would tell us", not "when the sequencer says it happened" — an upper bound. Both are recorded so the gap is visible.

7 · Rules the code enforces

- **Never invent a number.** A field that cannot be measured is reported as unavailable. A submission that fails produces a row with a null hash, never a placeholder.
- **Never publish a cost without its assumptions.** The ETH rate, its source, and the L1 gas price appear beside every figure.
- **Never report a latency without its finality level.**
- **Raw per-transaction data is primary.** Summaries are derived and regenerable; they never replace the rows.
- **Absolute timestamps only.** No duration is stored at collection time.
- **Observed and derived values are distinguished,** because a computed challenge-window deadline is not a witnessed event.
- **Do not stress-test a shared public network.** Achieved throughput is reported; no peak figure is claimed.

8 · What comes next

TASKS	WORK	WHY IT MATTERS NOW
E1–E3	Optimistic rollup adapter and cross-architecture run	The project plan names Optimistic Rollups explicitly, and the ZK-versus-optimistic inversion is the finding none of the surveyed papers produced. OP Sepolia is funded with its token deployed.
F2 · F3	Repetitions and batch-size sweep	Turns settlement from $n = 1$ into a distribution. Should start immediately and run in the background — it costs waiting, not work.
F4	Read-only mainnet observation	Testnet proof cadence is not mainnet's. Costs nothing and grounds every testnet figure.
G1–G5	Analysis and threats to validity	Several limitations are already documented here and belong in that section verbatim.
H1	Framework comparison table	No code, no data, no funds. Four of its six source papers are already in the project directory.

Outstanding decisions that need the supervisor

The project plan refers to a *"modular research testbed already designed and developed in our laboratory for ZK rollup technology"*. This implementation is deliberately observational and hosts nothing. Those are different approaches, and the difference has a concrete consequence: the plan asks for evaluation *"under varying load conditions"*, but a shared public sequencer must not be saturated — doing so would degrade a service other people depend on and would measure our own rate limits rather than the system. On public

testnets the only load variable under our control is batch size. A lab testbed would remove that constraint. This is worth settling early rather than in Week 10.